

# Operating Systems

## Filesystems

---

Sven Dziadek

October/November 2022


EPITA



# Storage devices

# Recall: Storage Types

Today, we are talking about **secondary storage**:



# Storage devices

---

Also called **mass storage**

Through the ages, lots of different norms & physical connections

The main ones:

- ATA
- Parallel ATA (IDE)
- SCSI
- ATAPI
- SATA
- NVMe also over PCIe


# Hard Disk Drive

## Seek time

Time to move the disk arm  
to desired cylinder

Sequential access preferred  
over random access

**Note:** Try storing data sequentially



# Non-Volatile Memory (NVM)

## Examples

SSD, flash memory storage (smartphone)

Access of data is **electronically** compared to hdd (electromechanically)

⇒ fast random access performance

Data cannot be overwritten

⇒ Copy changed data to new cell and mark old cell as unused

## Garbage collection

Larger units of unused NAND cells are erased

**Note:** Tell SSD controller when deleting files (TRIM)

## Some Similarities

### Block Devices

Divided into blocks of fixed size

For all operations: specify the Logical Block Address (LBA)

### Controllers

Kernel never talks with the drives directly

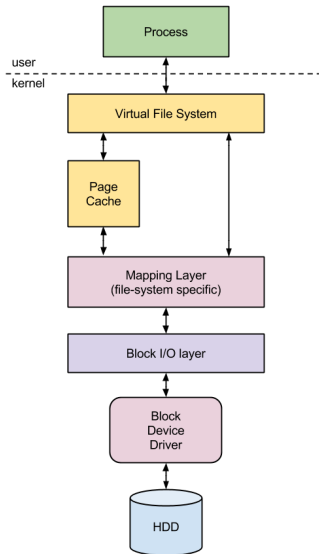
Give orders to controller

# Kernel Abstraction

## Block layer

Abstraction

to interact with block devices generically





# **Filesystems concepts**

## Filesystem

Method and data structure for secondary storage

**Lots** of different filesystems.

- Each has different features, purposes, limits and caveats, target OS.

Organization in the form of files & directories,  
features like Access-Control-Lists (ACLs) (for permissions) etc. . .


# Partitions

Disk partitions allow multiple filesystems

Depending on machine's firmware (BIOS / EFI)

Contain:

- location (start LBA)
- size (in disk blocks)
- (filesystem) type
- etc



# Partitions layout - Master Boot Record (MBR)


## Master Boot Record (MBR)

Legacy boot sector format

- Uses first sector of the disk (512 bytes)
- Loaded by the BIOS
- Boot routine (bootloader)
- 4 primary partitions table entries
- Partitions maximum size of 2 TiB ( $2^{32} \times 512$  bytes)

### Extended partitions:

pointer to another partition table



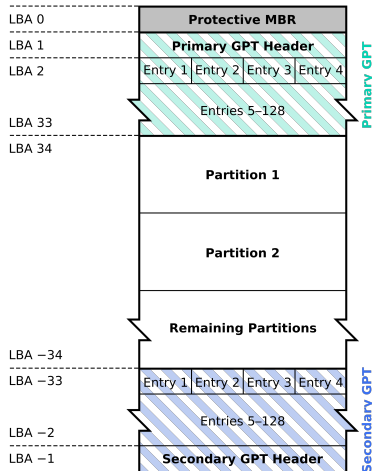
# Partitions layout - GUID Partition Table (GPT)

## GUID Partition Table (GPT)

Modern boot sector format

- Used with EFI firmwares
- Up to 128 entries
- Duplicated tables for more robustness
- Can boot full binaries  
in **EFI System Partitions (ESP)**
- Partition maximum size of 8 ZiB

## GUID Partition Table Scheme



## Superblock

Metadata of the filesystem

- Contains, e.g.,
  - Size
  - Block size
  - Number of free blocks
  - State
  - Root directory entry

**Note:** Corruption of the superblock makes whole filesystem unreadable  
⇒ usually **duplicated** a few times on the disk

# Filesystem - Inodes

Linux: Inode

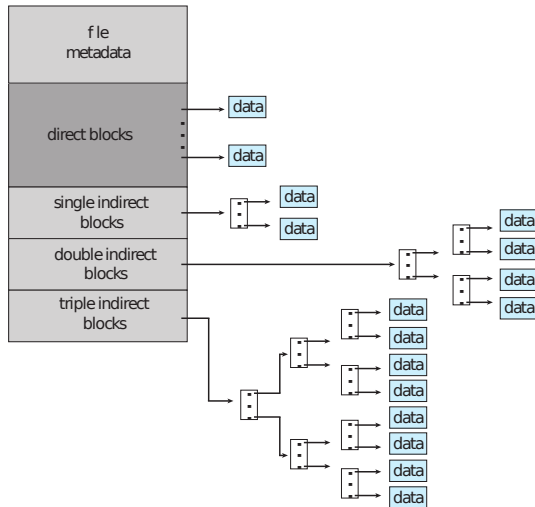
also called **file-control block**

Metadata of a file:

- Access, write, create time
- ACLs, owner, group
- File size
- More depending on the filesystem

Space for inodes is  
pre-allocated on filesystem creation

**Note:** Inode does **not** contain file names!



## Directories

Relate file names to inode

Main solutions:

**Linear list**

Simple and small

Linear search is slow

**Hash table**

Fast

Not easily adjust to size  
increase

## Hard link

Multiple directory entries pointing to one inode

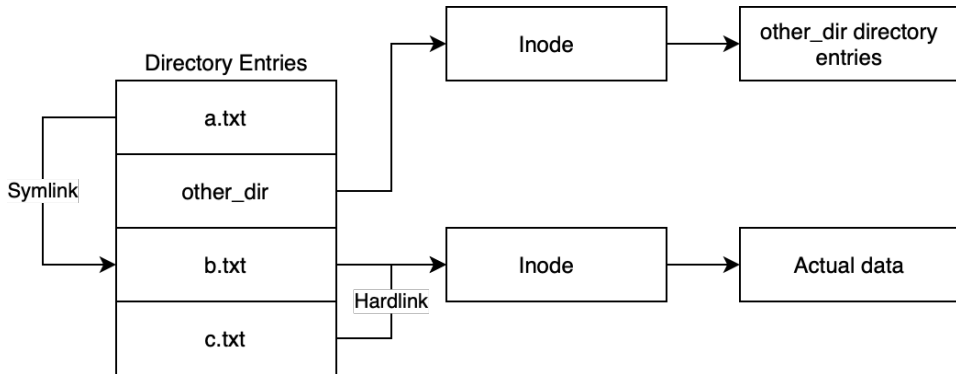
Unix typically stores also type in directory entry

Directories are files

⇒ they have an inode, directory entries are in file data



# Filesystem - Summary



## Demo: Inodes

---

Use `ln` or `ln -s` to create hard or symbolic links.

Compare their inodes with `ls -il` and `stat`.

Also compare inode of a directory with the inode of `.` in this directory.

# Free-Space List


Free blocks can be kept in **bit vector**  
but slow for large disks

**Linked list** give first free block fast

For larger allocation:

**grouping** (link to  $n$  free blocks)


**counting** (add counter to links if multiple  
contiguous blocks are free)



# Caching

Keep important parts in memory:

- currently used directories
- free-space list
- used blocks
- for sequential access:  
subsequent blocks



After a crash, filesystem must be checked for inconsistency

To simplify and speed up recovery: use log

- all metadata changes are written to log
- replay log in real fs structures
- mark changes in log as done
- After system crash: complete unfinished transactions

Some filesystems like ZFS or btrfs go into volume management territory with:

- Logical volumes
- Checksumming
- Copy-on-write & snapshotting
- Software RAID

**Unix way of life**

# Virtual Filesystem (VFS)

**Everything is a file**

Many things abstracted as file-like objects

Handled as files:

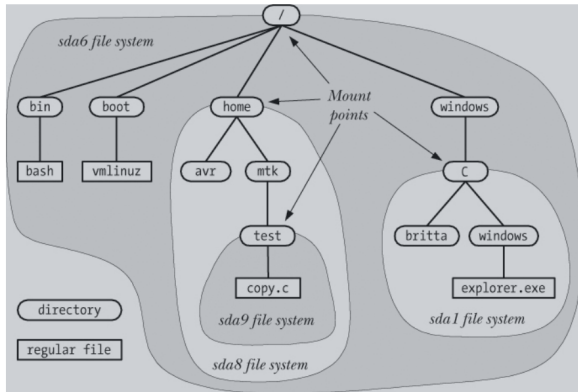
- Actual files
- Unix Sockets, Named pipes
- Devices (**char devices** & **block devices**)
- Partitions
- Kernel interfaces (sysfs and procfs)



# Virtual Filesystem (VFS)

Direct interface to user applications:  
Filesystem implementation  
completely abstracted

Filesystems are **mounted** on the VFS  
(e.g your **rootfs** is mount on /)



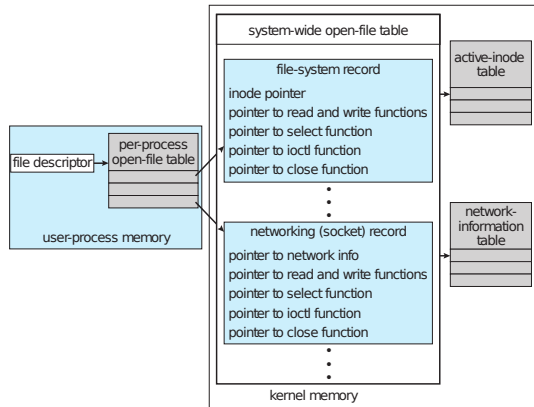
```
cat /proc/self/status
```

For more info:

<https://www.kernel.org/doc/html/latest/filesystems/proc.html>

# File descriptors

- Kernel keeps a table of all opened file descriptors
- Each process has its own table: the **File Descriptor Table**
- Entries of those tables represent the file, and its **operations**  
Syscalls `read(2)`, `seek(2)`, ... are linked to those operations
- `open(2)`, `socket(2)`, `pipe(2)`, ... create new entries
- The file descriptor integers in user space are the **index in the table**



# Device mapper

---

- Devices in the VFS are **logical mappings** to physical devices
- This is managed through the **device mapper** subsystem
- This gives a lot of flexibility on how you manage block devices
- Foundation for other subsystems:
  - LVM: for logical volumes management
  - `dm-crypt`: for disk encryption
  - Software RAID

## Special files - ioctl(2)

- Some **files** in the VFS are special files - e.g **char devices**
- They can have special operations
  - For example, changing the baudrate of a serial tty
- Those operations are done through the ioctl(2) syscall
- A bit of a **do anything** kind of syscall
- The available commands depend on the file and the underlying driver

Please take the time to fill the following form:

<https://forms.forge.epita.fr/>

Thank you for your feedback!

