

Operating Systems

Threads

Sven Dziadek

October/November 2022

EPITA



Pthreads

POSIX Threads: Pthreads

Link with `-pthread`

```
#include <pthread.h>
```

Create threads

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*start_routine) (void *), void *arg);  
int pthread_join(pthread_t thread, void **retval);
```

Race Conditions

Race Conditions

Race Condition

Result depends on timing

If some possible results are undesirable \Rightarrow bug

Example (Too-much-milk problem)

Consider the following algorithm in a shared flat:

- Check fridge for milk
- Go to buy milk
- Put milk into fridge



Race Condition: Demo

Start **two threads** that increment a shared variable `c`:

```
for(i=0;i<1000;i++)  
    c++
```

What is the final value of `c`?

See `race.c`

Race Condition: Demo

c++ translates to:

```
0x11a3 <fnC+30>:  mov    0x2eb3(%rip),%eax    # load 0x405c <c> into %eax
0x11a9 <fnC+36>:  add     $0x1,%eax                # %eax = %eax + 1
0x11ac <fnC+39>:  mov     %eax,0x2eaa(%rip)         # store %eax to 0x405c <c>
```

Thread 1

load c =0

increment c =1

store c =1

Thread 2

load c =0

increment c =1

store c =1

Synchronization Tools

Hardware: Atomic Operations

Atomic Operations in Hardware

CPUs can block memory bus for a few instructions

GCC provides (for any integral scalar or pointer type as type):

```
type __atomic_add_fetch (type *ptr, type val, int memorder);  
type __atomic_fetch_add (type *ptr, type val, int memorder);  
bool __atomic_compare_exchange (type *ptr, type *expected, type *desired, bool weak,  
                                int success_memorder, int failure_memorder);
```

Also for more operations

See https://gcc.gnu.org/onlinedocs/gcc-10.2.0/gcc/_005f_005fatomic-Builtins.html

Atomic Operations: Demo

race.c

Pthreads: Mutex

Lock for **mutual exclusion**:

- lock()
- critical section...
- unlock()

```
pthread_mutex_t fastmutex = PTHREAD_MUTEX_INITIALIZER;  
int pthread_mutex_destroy(pthread_mutex_t *mutex);  
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_trylock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

Mutex Demo

race.c

Locking: Notions

acquire lock

critical section

release lock

Critical Section

Block of code (accessing shared resources)
that must only be entered **at most once** at a time

Effective goal: **serialization**

Contention

Try to acquire a lock that is already locked

Acquire a free lock = uncontended

Note: High contention decreases performance

⇒ keep critical sections short / use **finegrained** locking

Pthreads: Spinlock

Like a mutex but

does not block

To be used for short duration only

(less than time of two context switches)

with real-time scheduling policies

```
int pthread_spin_init(pthread_spinlock_t *lock, int pshared);  
int pthread_spin_destroy(pthread_spinlock_t *lock);  
int pthread_spin_lock(pthread_spinlock_t *lock);  
int pthread_spin_trylock(pthread_spinlock_t *lock);  
int pthread_spin_unlock(pthread_spinlock_t *lock);
```

Pthreads: Semaphores

```
#include <semaphore.h>
int sem_init(sem_t *sem, int pshared, unsigned int value);
int sem_destroy(sem_t *sem);
int sem_wait(sem_t *sem);
int sem_post(sem_t *sem);
sem_t *sem_open(const char *name, int oflag, mode_t mode, unsigned int value);
```

Value indicates how many resources are free

Blocks if value would become negative

See `sem_overview(7)`

Monitors

Monitor

Construct or abstract data type

- for mutual exclusion
- with ability to wait for conditions

Example: Java's **synchronized** methods

Condition variables allow to

- wait for conditions to be fulfilled
- block on false conditions
and release of lock while waiting

Pattern to be solved:

```
lock()
while (!condition) { //must be atomic
    unlock()
    sleep() //release lock while waiting
    lock()
}
/* critical section */
unlock()
```

Monitor Example: Producer-Consumer Problem

Producer-Consumer Problem

Resource synchronization problem
to distribute work packages/data

Producer threads distribute work packages

Consumer threads collect work packages

Assume a ring buffer of fixed size

Producer threads can push to buffer
if not full

Consumer threads can pop from buffer
if not empty

Note:

Check for empty/full buffer
only while holding **lock**

If condition not fulfilled
⇒ release lock while waiting
+ blocking while waiting preferred

Solution:

Use **mutex** together
with **condition variables**

Pthreads: Condition Variables

```
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;  
int pthread_cond_destroy(pthread_cond_t *cond);  
int pthread_cond_signal(pthread_cond_t *cond);  
int pthread_cond_broadcast(pthread_cond_t *cond);  
int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);
```

Note that wait unlocks mutex while waiting

Use signal or broadcast depending on how many threads should be woken up


Deadlocks

Deadlock

Deadlock

Threads are blocking each other

Must be prevented at all costs




Note: Locking mutex twice can also lead to deadlock (Pthreads allows recursive mutex)

Deadlock: Necessary Conditions

- Mutual Exclusion
- Resources cannot be preempted
- Hold one and wait for other resource
- Circular wait

Dining Philosopher's Problem



Algorithm for each of the five philosophers:

- Think and then take left chopstick
- Think and then take right chopstick
- Eat some rice
- Put left chopstick back
- Put right chopstick back
- Repeat

Possible solutions: Break symmetry or atomically pick both chopsticks

Linux Kernel Support

Fast Userspace Mutex

System call `futex(2)` only called in contended case

Fast in uncontended case (only atomic operation)

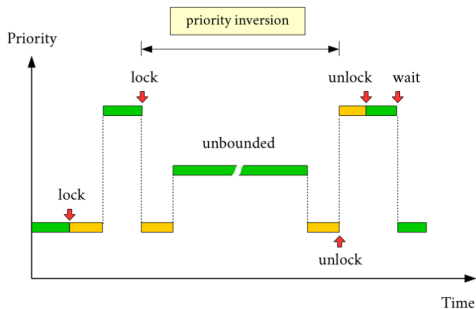
Contended case: kernel manages wait queue

Not normally used directly
but used in Pthread library

See `futex(7)`

More Theory

Priority Inversion in Real-Time Scheduling



Priority Inversion (PI)

Blocking of a high-priority job
by a lower-priority job

Unbounded PI must be prevented

Peterson's Algorithm (For two threads – thread 0 and thread 1)

```
int turn = 0; char flag[2] = {0, 0};
```

```
void lock(int me) {  
    int other = 1-me;
```

```
    flag[me] = 1;
```

```
    turn = other;
```

```
    while (flag[other]  
           && turn == other)  
        ; //busy wait
```

```
}
```

```
void unlock(int me) {
```

```
    flag[me] = 0;
```

```
}
```

Peterson's Algorithm

Provably correct for mutual exclusion

Guarantees progress and bounded waiting

```
/* I want to enter */
```

```
/* You can enter next */
```

```
/* If you want to enter */
```

```
/* and it's your turn */
```

```
/* I'll wait */
```

```
/* I don't want to enter anymore */
```

Why not to build your own synchronization tools

Demo: Peterson's Algorithm does not work on modern machines

Compiler & Processor Optimizations

Compiler can reorder instructions

CPU reorders memory accesses

Memory barriers can help: `__atomic_thread_fence(__ATOMIC_SEQ_CST)`

Demo: Peterson's Algorithm works with memory barrier

Note: Writing synchronization tools is hard \Rightarrow use existing tools!